

task-vit

April 5, 2024

```
[2]: import os
import cv2
import h5py
import pyarrow.parquet as pq
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from tqdm import tqdm
from statistics import mean
from sklearn.metrics import accuracy_score, roc_auc_score
from sklearn.model_selection import StratifiedKFold, train_test_split
from sklearn.preprocessing import StandardScaler

import torch
import torch.nn as nn
import torch.optim as optim
from torch.optim import lr_scheduler
from torchlars import LARS
from sklearn.metrics import roc_auc_score
from transformers import AutoModel, AutoTokenizer
import torchmetrics
from torch.utils.data import DataLoader, TensorDataset
from torch.optim.lr_scheduler import CosineAnnealingLR, StepLR, ReduceLROnPlateau
from torch.utils.data.sampler import SubsetRandomSampler, BatchSampler, Sampler
from torch.optim import Adam, SGD
from torchvision import transforms, models
from torch.utils.data import Dataset
from torchvision.transforms import Resize, ToTensor
from torch.cuda.amp import autocast, GradScaler
import torch.nn.functional as F
from torch.nn.utils import clip_grad_norm_

torch.manual_seed(42)
np.random.seed(42)
torch.cuda.manual_seed(42)
```

```
[3]: gpus = torch.cuda.device_count()
    if gpus <= 1:
        device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
        print(f'Using {gpus} GPU')
    else:
        device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
        print(f'Using {gpus} GPUs')
```

Using 2 GPUs

```
[4]: MIX = False
    if MIX:
        scaler = GradScaler()
        print('Mixed precision enabled')
    else:
        print('Using full precision')
```

Using full precision

```
[5]: columns_to_keep = list(range(22))
    dataset = pd.read_csv("/kaggle/input/vit-data/HIGGS.csv/HIGGS.csv",
        ↪usecols=columns_to_keep)
```

```
[6]: dataset.head()
```

```
[6]: 1.0000000000000000e+00  8.692932128906250000e-01  \
0          1.0          0.907542
1          1.0          0.798835
2          0.0          1.344385
3          1.0          1.105009
4          0.0          1.595839

-6.350818276405334473e-01  2.256902605295181274e-01  \
0          0.329147          0.359412
1          1.470639         -1.635975
2         -0.876626          0.935913
3          0.321356          1.522401
4         -0.607811          0.007075

3.274700641632080078e-01 -6.899932026863098145e-01  \
0          1.497970         -0.313010
1          0.453773          0.425629
2          1.992050          0.882454
3          0.882808         -1.205349
4          1.818450         -0.111906

7.542022466659545898e-01 -2.485731393098831177e-01  \
0          1.095531         -0.557525
```

1	1.104875	1.282322	
2	1.786066	-1.646778	
3	0.681466	-1.070464	
4	0.847550	-0.566437	
	-1.092063903808593750e+00	0.000000000000000000e+00	... \
0	-1.588230	2.173076	...
1	1.381664	0.000000	...
2	-0.942383	0.000000	...
3	-0.921871	0.000000	...
4	1.581239	2.173076	...
	9.303491115570068359e-01	1.107436060905456543e+00	\
0	1.271015	2.214872	
1	-0.819690	2.214872	
2	0.736159	2.214872	
3	0.971407	2.214872	
4	1.426367	0.000000	
	1.138904333114624023e+00	-1.578198313713073730e+00	\
0	0.499994	-1.261432	
1	0.993490	0.356080	
2	1.298720	-1.430738	
3	0.596761	-0.350273	
4	0.921661	-1.190432	
	-1.046985387802124023e+00	0.000000000000000000e+00.1	\
0	0.732156	0.000000	
1	-0.208778	2.548224	
2	-0.364658	0.000000	
3	0.631194	0.000000	
4	-1.615589	0.000000	
	6.579295396804809570e-01	-1.045456994324922562e-02	\
0	0.398701	-1.138930	
1	1.256955	1.128848	
2	0.745313	-0.678379	
3	0.479999	-0.373566	
4	0.651114	-0.654227	
	-4.576716944575309753e-02	3.101961374282836914e+00	
0	-0.000819	0.000000	
1	0.900461	0.000000	
2	-1.360356	0.000000	
3	0.113041	0.000000	
4	-1.274345	3.101961	

[5 rows x 22 columns]

```
[7]: dataset.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10999999 entries, 0 to 10999998
Data columns (total 22 columns):
 #   Column                                Dtype
---  -
 0   1.0000000000000000e+00                float64
 1   8.692932128906250000e-01              float64
 2   -6.350818276405334473e-01            float64
 3   2.256902605295181274e-01            float64
 4   3.274700641632080078e-01            float64
 5   -6.899932026863098145e-01            float64
 6   7.542022466659545898e-01            float64
 7   -2.485731393098831177e-01            float64
 8   -1.092063903808593750e+00            float64
 9   0.0000000000000000e+00              float64
10   1.374992132186889648e+00            float64
11   -6.536741852760314941e-01            float64
12   9.303491115570068359e-01            float64
13   1.107436060905456543e+00            float64
14   1.138904333114624023e+00            float64
15   -1.578198313713073730e+00            float64
16   -1.046985387802124023e+00            float64
17   0.0000000000000000e+00.1            float64
18   6.579295396804809570e-01            float64
19   -1.045456994324922562e-02            float64
20   -4.576716944575309753e-02            float64
21   3.101961374282836914e+00            float64
dtypes: float64(22)
memory usage: 1.8 GB
```

```
[8]: y = dataset.iloc[:,0]
     X = dataset.iloc[:, 1:]
```

```
[9]: class_distribution = y.value_counts()
     print("Class Distribution:")
     print(class_distribution)
```

```
Class Distribution:
1.0000000000000000e+00
1.0      5829122
0.0      5170877
Name: count, dtype: int64
```

```
[10]: X = (X - np.mean(X))/np.std(X)
```

```

/opt/conda/lib/python3.10/site-packages/numpy/core/fromnumeric.py:3643:
FutureWarning: The behavior of DataFrame.std with axis=None is deprecated, in a
future version this will reduce over both axes and return a scalar. To retain
the old behavior, pass axis=0 (or do not pass axis)
    return std(axis=axis, dtype=dtype, out=out, ddof=ddof, **kwargs)

```

```

[11]: X_train, X_val, y_train, y_val = train_test_split(X, y, test_size=100000,
↳stratify = y, random_state=42)

```

```

[12]: X_train = torch.tensor(X_train.values, dtype=torch.float32)
X_val = torch.tensor(X_val.values, dtype=torch.float32)
y_train = torch.tensor(y_train.values, dtype=torch.float32)
y_val = torch.tensor(y_val.values, dtype=torch.float32)

```

```

[13]: class CFG:
    lr = 1e-3
    weight_decay = 5e-4
    input_dim = 21
    latent_dim = 64
    num_classes = 1
    num_layers = 3
    num_heads = 3
    dropout = 0.1
    num_epochs = 17
    batch_size = 1024

```

```

[14]: train_dataset = TensorDataset(X_train, y_train)
val_dataset = TensorDataset(X_val, y_val)
train_loader = DataLoader(train_dataset, batch_size=CFG.batch_size,
↳shuffle=True)
val_loader = DataLoader(val_dataset, batch_size=CFG.batch_size)

```

```

[15]: class TransformerAutoencoderClassifier(nn.Module):
    def __init__(self, input_dim=21, latent_dim=64, num_classes=1,
↳num_layers=2, num_heads=3, dropout=0.1):
        super(TransformerAutoencoderClassifier, self).__init__()

        self.input_dim = input_dim
        self.num_layers = num_layers
        self.num_heads = num_heads

        self.encoder_layers = nn.TransformerEncoderLayer(d_model=input_dim,
↳nhead=num_heads)
        self.encoder = nn.TransformerEncoder(self.encoder_layers,
↳num_layers=num_layers)
        self.decoder_layers = nn.TransformerDecoderLayer(d_model=input_dim,
↳nhead=num_heads)

```

```

        self.decoder = nn.TransformerDecoder(self.decoder_layers,
↪num_layers=num_layers)
        self.bottleneck = nn.Linear(input_dim, latent_dim)
        self.classifier = nn.Linear(latent_dim, num_classes)
        self.dropout = nn.Dropout(dropout)

    def forward(self, x):

        tgt_mask = self.generate_square_subsequent_mask(x.size(0))
        encoder_output = self.encoder(x)
        decoder_output = self.decoder(x, encoder_output, tgt_mask=tgt_mask)
        bottleneck_output = self.bottleneck(encoder_output)
        classification_output = torch.sigmoid(self.
↪classifier(bottleneck_output))
        return decoder_output, classification_output

    def generate_square_subsequent_mask(self, sz):
        mask = (torch.triu(torch.ones(sz, sz)) == 1).transpose(0, 1)
        mask = mask.float().masked_fill(mask == 0, float('-inf')).
↪masked_fill(mask == 1, float(0.0))
        return mask

```

```

[16]: model = TransformerAutoencoderClassifier(
        input_dim=CFG.input_dim,
        latent_dim=CFG.latent_dim,
        num_classes=CFG.num_classes,
        num_layers=CFG.num_layers,
        num_heads=CFG.num_heads,
        dropout=CFG.dropout
    ).to(device)

```

/opt/conda/lib/python3.10/site-packages/torch/nn/modules/transformer.py:282: UserWarning: enable_nested_tensor is True, but self.use_nested_tensor is False because encoder_layer.self_attn.batch_first was not True(use batch_first for better inference performance)

warnings.warn(f"enable_nested_tensor is True, but self.use_nested_tensor is False because {why_not_sparsity_fast_path}")

```

[17]: criterion_autoencoder = nn.MSELoss()
        criterion_classification = nn.BCELoss()
        optimizer = torch.optim.AdamW(model.parameters(), lr=CFG.lr, weight_decay = CFG.
↪weight_decay)
        scheduler = ReduceLROnPlateau(optimizer, mode='min', patience=1, factor=0.4,
↪threshold=1e-2, verbose=True)

```

```

[18]: model.eval()
        val_running_loss = 0

```

```

correct_val = 0
total_val = 0
y_true = []
y_scores = []

with torch.no_grad():
    with tqdm(val_loader, desc="Validation", leave=False) as val_loader_with_progress:
        for val_inputs, val_labels in val_loader_with_progress:
            val_inputs, val_labels = val_inputs.to(device), val_labels.to(device)

            decoder_output, classification_output = model(val_inputs)
            autoencoder_loss = criterion_autoencoder(decoder_output, val_inputs)

            classification_loss = criterion_classification(classification_output.squeeze(), val_labels.float())
            outputs = classification_output.squeeze()
            val_loss = 0.5 * autoencoder_loss + 0.5 * classification_loss
            val_running_loss += val_loss.item()
            predicted_val = torch.round(outputs)
            correct_val += (predicted_val == val_labels).sum().item()
            total_val += val_labels.size(0)
            y_true += val_labels.cpu().detach().numpy().tolist()
            y_scores += outputs.cpu().detach().numpy().tolist()
            val_loader_with_progress.set_postfix(val_loss=val_loss.item())

val_loss = val_running_loss / len(val_loader)
val_acc = 100 * correct_val / total_val
val_roc_auc = roc_auc_score(y_true, y_scores)

print(f'At 0th epoch without training Validation Loss: {val_loss}, Validation AUC-ROC: {val_roc_auc}, Validation Accuracy: {val_acc}')

```

At 0th epoch without training Validation Loss: 0.9273783120573783, Validation AUC-ROC: 0.5367654452475378, Validation Accuracy: 53.759

```

[19]: train_losses = []
      val_losses = []
      train_accs = []
      val_accs = []
      train_roc_aucs = []
      val_roc_aucs = []
      best_val_accuracy = 0.0
      best_model_state_dict = None

      for epoch in range(CFG.num_epochs):

```

```

model.train()
train_running_loss = 0
correct_train = 0
total_train = 0
y_true = []
y_scores = []

with tqdm(train_loader, desc=f"Epoch {epoch + 1}/{CFG.num_epochs}",
↪leave=False) as train_loader_with_progress:
    for inputs, labels in train_loader_with_progress:
        inputs, labels = inputs.to(device), labels.to(device)
        optimizer.zero_grad()
        decoder_output, classification_output = model(inputs)
        autoencoder_loss = criterion_autoencoder(decoder_output, inputs)
        classification_loss = ↪
↪criterion_classification(classification_output.squeeze(), labels.float())
        loss = 0.5 * autoencoder_loss + 0.5 * classification_loss
        outputs = classification_output.squeeze()
        loss.backward()
        clip_grad_norm_(model.parameters(), max_norm=3)
        optimizer.step()
        train_running_loss += loss.item()
        predicted_train = torch.round(outputs)
        correct_train += (predicted_train == labels).sum().item()
        total_train += labels.size(0)
        train_loader_with_progress.set_postfix(train_loss=loss.item())
        y_true += labels.cpu().detach().numpy().tolist()
        y_scores += outputs.cpu().detach().numpy().tolist()

train_loss = train_running_loss / len(train_loader)
train_acc = 100 * correct_train / total_train
roc_auc_train = roc_auc_score(y_true, y_scores)

train_losses.append(train_loss)
train_accs.append(train_acc)
train_roc_aucs.append(roc_auc_train)

model.eval()
val_running_loss = 0
correct_val = 0
total_val = 0
y_true = []
y_scores = []

with torch.no_grad():
    with tqdm(val_loader, desc=f"Epoch {epoch + 1}/{CFG.num_epochs}",
↪Validation", leave=False) as val_loader_with_progress:

```



```

        for val_inputs, val_labels in val_loader_with_progress:
            val_inputs, val_labels = val_inputs.to(device), val_labels.
            to(device)

            decoder_output, classification_output = model(val_inputs)
            autoencoder_loss = criterion_autoencoder(decoder_output,
            val_inputs)

            classification_loss =
            criterion_classification(classification_output.squeeze(), val_labels.float())
            outputs = classification_output.squeeze()
            val_loss = 0.5 * autoencoder_loss + 0.5 * classification_loss
            val_running_loss += val_loss.item()
            predicted_val = torch.round(outputs)
            correct_val += (predicted_val == val_labels).sum().item()
            total_val += val_labels.size(0)
            y_true += val_labels.cpu().detach().numpy().tolist()
            y_scores += outputs.cpu().detach().numpy().tolist()
            val_loader_with_progress.set_postfix(val_loss=val_loss.item())

    val_loss = val_running_loss / len(val_loader)
    val_acc = 100 * correct_val / total_val
    val_roc_auc = roc_auc_score(y_true, y_scores)

    val_losses.append(val_loss)
    val_accs.append(val_acc)
    val_roc_aucs.append(val_roc_auc)
    scheduler.step(val_loss)
    for param_group in optimizer.param_groups:
        print(f'Learning Rate: {param_group["lr"]}')

    print(f"Epoch {epoch + 1}/{CFG.num_epochs}, Train Loss: {train_loss:.4f},
    Train Acc: {train_acc:.2f}%, Val Loss: {val_loss:.4f}, Val Acc: {val_acc:.
    2f}%, Train ROC-AUC: {roc_auc_train:.3f}, Val ROC-AUC: {val_roc_auc:.3f}")

    if val_acc > best_val_accuracy:
        best_val_accuracy = val_acc
        best_model_state_dict = model.state_dict()
        torch.save(best_model_state_dict,
        "model_weights_Transformer_Autoencoder.pth")

print("Model saved successfully.")
print("Finished training")

```

Learning Rate: 0.001

Epoch 1/17, Train Loss: 0.3187, Train Acc: 67.89%, Val Loss: 0.2932, Val Acc: 70.56%, Train ROC-AUC: 0.743, Val ROC-AUC: 0.777

Learning Rate: 0.001

Epoch 2/17, Train Loss: 0.2949, Train Acc: 70.38%, Val Loss: 0.2836, Val Acc: 71.82%, Train ROC-AUC: 0.773, Val ROC-AUC: 0.794

Learning Rate: 0.001

Epoch 3/17, Train Loss: 0.2893, Train Acc: 71.23%, Val Loss: 0.2786, Val Acc: 72.73%, Train ROC-AUC: 0.783, Val ROC-AUC: 0.802

Learning Rate: 0.001

Epoch 4/17, Train Loss: 0.2856, Train Acc: 71.77%, Val Loss: 0.2750, Val Acc: 73.20%, Train ROC-AUC: 0.789, Val ROC-AUC: 0.808

Learning Rate: 0.001

Epoch 5/17, Train Loss: 0.2831, Train Acc: 72.11%, Val Loss: 0.2722, Val Acc: 73.42%, Train ROC-AUC: 0.794, Val ROC-AUC: 0.812

Learning Rate: 0.001

Epoch 6/17, Train Loss: 0.2814, Train Acc: 72.35%, Val Loss: 0.2699, Val Acc: 73.83%, Train ROC-AUC: 0.797, Val ROC-AUC: 0.816

Learning Rate: 0.001

Epoch 7/17, Train Loss: 0.2800, Train Acc: 72.57%, Val Loss: 0.2692, Val Acc: 74.04%, Train ROC-AUC: 0.799, Val ROC-AUC: 0.818

Learning Rate: 0.001

Epoch 8/17, Train Loss: 0.2790, Train Acc: 72.71%, Val Loss: 0.2682, Val Acc: 74.14%, Train ROC-AUC: 0.801, Val ROC-AUC: 0.819

Epoch 00009: reducing learning rate of group 0 to 4.0000e-04.

Learning Rate: 0.0004

Epoch 9/17, Train Loss: 0.2779, Train Acc: 72.86%, Val Loss: 0.2672, Val Acc: 74.13%, Train ROC-AUC: 0.803, Val ROC-AUC: 0.820

Learning Rate: 0.0004

Epoch 10/17, Train Loss: 0.2743, Train Acc: 73.35%, Val Loss: 0.2626, Val Acc: 74.84%, Train ROC-AUC: 0.809, Val ROC-AUC: 0.828

Learning Rate: 0.0004

Epoch 11/17, Train Loss: 0.2736, Train Acc: 73.43%, Val Loss: 0.2623, Val Acc: 74.93%, Train ROC-AUC: 0.810, Val ROC-AUC: 0.828

Epoch 00012: reducing learning rate of group 0 to 1.6000e-04.

Learning Rate: 0.00016

Epoch 12/17, Train Loss: 0.2731, Train Acc: 73.50%, Val Loss: 0.2620, Val Acc: 75.00%, Train ROC-AUC: 0.811, Val ROC-AUC: 0.829

Learning Rate: 0.00016

Epoch 13/17, Train Loss: 0.2716, Train Acc: 73.70%, Val Loss: 0.2602, Val Acc: 75.25%, Train ROC-AUC: 0.813, Val ROC-AUC: 0.831

Epoch 00014: reducing learning rate of group 0 to 6.4000e-05.

Learning Rate: 6.400000000000001e-05

Epoch 14/17, Train Loss: 0.2712, Train Acc: 73.75%, Val Loss: 0.2604, Val Acc: 75.17%, Train ROC-AUC: 0.814, Val ROC-AUC: 0.831

Learning Rate: 6.400000000000001e-05

Epoch 15/17, Train Loss: 0.2706, Train Acc: 73.84%, Val Loss: 0.2595, Val Acc: 75.28%, Train ROC-AUC: 0.815, Val ROC-AUC: 0.832

Learning Rate: 6.400000000000001e-05

Epoch 16/17, Train Loss: 0.2705, Train Acc: 73.85%, Val Loss: 0.2594, Val Acc: 75.30%, Train ROC-AUC: 0.815, Val ROC-AUC: 0.833

Epoch 00017: reducing learning rate of group 0 to 2.5600e-05.


Learning Rate: 2.5600000000000006e-05

Epoch 17/17, Train Loss: 0.2704, Train Acc: 73.86%, Val Loss: 0.2594, Val Acc: 75.26%, Train ROC-AUC: 0.815, Val ROC-AUC: 0.833


Model saved successfully.

Finished training

```
[20]: plt.plot(train_losses, label="Training Loss")
plt.plot(val_losses, label="Val Loss")
plt.xlabel("Epoch")
plt.ylabel("Loss")
plt.title("Training and Val Loss")
plt.legend()
plt.show()
```



```
[21]: plt.plot(train_accs, label="Training Accuracy")
plt.plot(val_accs, label="Val Accuracy")
plt.xlabel("Epoch")
plt.ylabel("Accuracy (%)")
plt.title("Training and Val Accuracy")
plt.legend()
plt.show()
```



```
[22]: plt.plot(train_roc_aucs, label="Training ROC")
plt.plot(val_roc_aucs, label="Val ROC")
plt.xlabel("Epoch")
plt.ylabel("ROC-AUC")
plt.title("Training and Val ROC-AUC")
plt.legend()
plt.show()
```

